Does Remote Sensing Beat the Trail Map? Predicting Hiking Speed in Unseen Regions from Lidar Terrain and OpenStreetMap Tags

Brian Zhou [TODO: author list and order: ask Danni, Tianyi and Prof. Jacobs]
Department of Computer Science and Engineering, Washington University in St. Louis

***Abstract*—[TODO: 150–200 words, written last. Problem: travel time on foot; Tobler. Data: 17 regions, 290k walking segments from public OSM GPS traces. Test: leave one region out, calibrated predictive distributions (CRPS). Result: terrain from airborne-lidar elevation transfers to unseen regions (+2.1 to +2.8% CRPS skill over a refitted slope curve), map tags barely do (+0.4 to +0.9%), satellite embeddings help within regions but transfer less reliably; Tobler is 12.5% worse. A terrain-adjusted hiking function needs one DEM-derived input.]**

***Index Terms*—hiking speed, lidar, OpenStreetMap, GPS trajectories, transfer, calibration**

I. Introduction

[TODO: Paragraph 1, why travel time on foot matters: search and rescue, wildfire escape routes, trail apps, accessibility maps. GIS still uses Tobler's 1993 hiking function [1], $v = 6\exp(-3.5\,|s + 0.05|)$ km/h, and fixed multipliers for other modes.]

[TODO: Paragraph 2, what is known: crowdsourced GPS shows Tobler is too fast and too peaked [2], [3]; Wood et al. [4] add OSM terrain classes and lidar obstruction (England and Wales); Huang et al. [5] fit boosted trees with 16 factors, weather included (Hong Kong); Šerić et al. [6] predict a latent walking speed from Sentinel-2 (one Croatian range); HikingTTE [7] learns whole-hike time from the hiker's first part. STRIDE [8] models walking speed from lidar slope, understory density [9] and roughness, calibrated on field walks; RasterNet [10] learns road speed from lidar and imagery.]

[TODO: Paragraph 3, the gap: each of these studies fits and tests in one country, city or range with random or within-area splits (Wood: 10-fold cross-validation; Huang: random split); none tests whether terrain covariates transfer to regions a model has never seen, against the map tags used today, with calibrated uncertainty. Wood et al. also found no gain over existing functions for whole-route times (13.5–15.5% error for every model).]

[TODO: Paragraph 4, contributions (three bullets): (1) a 17-region dataset of walking traversals of 50 m trail segments from public OSM traces; (2) an evaluation protocol (leave one region out, CRPS, coverage) with published baselines; (3) the finding that elevation-derived terrain transfers and map tags do not, plus a one-input terrain-adjusted hiking function.]

II. Data

[TODO: Draft by Claude: rewrite in your own words.]

Sources. We use public GPS traces uploaded to OpenStreetMap (OSM) [11] in 17 mountain regions of the United States and Canada (Fig. 1), chosen for relief, trail density and lidar coverage. The traces come from three sources: the OSM trackpoints API, which returns the timestamped points of identifiable and trackable uploads tile by tile; the original GPX files of identifiable uploads, which give whole hikes at the device's own sampling; and a crawl of 1.5 million original files of public and identifiable uploads from 2016–2022, which keeps the timestamps the API drops. Sixteen regions combine all three; the Bay Area uses the crawl only, because the API times out on its tiles.

Processing. Each recording is cleaned (repeated timestamps, steps faster than 60 km/h), split at gaps over 5 min and thinned to points at least 5 s apart. A step counts as a stop when the point 60 s later lies within 10 m. Traces are map-matched to the OSM trail and road network with a hidden Markov model [12] (97% of points matched), and the matched path is cut into 50 m segments. Each segment records its moving speed, its signed slope along the direction of travel from the 5 m USGS 3DEP elevation grid [13] (resampled from 1 m lidar under at least 83% of walked trail in 13 regions), the OSM tags of its way and the trace's recording interval. Duplicate uploads are removed, and traces that start within 15 min and 1 km of each other form one outing (a group hike), the unit we resample. Travel mode is inferred per trace, anchored on trails closed to bikes and horses. Against the activity that 327 uploaders' apps wrote into their files, 93% of labelled hikes are classed as walking, and 5% of the labelled traces we class as walking say bike or ski.

The dataset. The regions hold 2,876 hiking traces. The walking analysis set (walking traces on trails) has 290,640 segments (14,532 km) from 2,046 traces in 1,992 outings and 4,702 moving hours. Its uploaders are OSM mappers and enthusiasts rather than a random sample of hikers, and busy API tiles do not reach their oldest uploads. Traces logged more than 10 s apart read up to 11% slow, because straight lines between points cut corners and short stops hide between them; the gains we report are the same at every recording rate (Sec. IV).

III. Methods

[TODO: Draft by Claude: rewrite in your own words.]

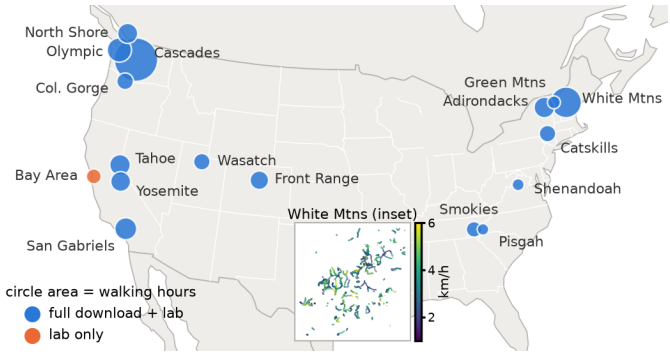


Fig. 1. The 17 regions. Circle area is walking hours (1,992 outings, 4,702 h in total). Blue: full OSM download plus the lab crawl; orange: lab crawl only. Inset: the White Mountains' 34,535 walking segments colored by speed.

Target and covariates. We predict the distribution of $\ln v$, the log moving speed of a walking segment. Covariates come in five groups: (a) slope; (b) OSM tags: sac_scale, trail_visibility, surface, width, smoothness, mtb_scale, highway type and permitted modes; (c) terrain from elevation, read every 1 m along the segment's matched path: the spread (SD) and maximum of its 5 m pitches, the relief around an 11 m moving average and the side slope, from the 5 m grid and, where available, 1 m lidar, which adds STRIDE's ground roughness [8]; (d) 2021 land cover and tree canopy at 30 m [14]; (e) 2022 AlphaEarth satellite embeddings at the segment midpoint [15], [16]. As a control, "hidden climbing" is the extra time Tobler's function predicts when the segment is walked as ten 5 m pieces instead of at its average grade.

Models. The baselines are Tobler's function [1], the published median curve of Campbell et al. [3], STRIDE [8] as published and with its form refitted, and our reference: a rounded, asymmetric slope curve refitted to the training data, $\ln v = a - k_{\pm}\left(\sqrt{(s+s_0)^2+e^2} - e\right)$, with k_+ above the peak and k_- below it. Each covariate group then enters two models: an additive one (ridge regression of the curve's residuals on one-hot bins: deciles of continuous covariates, one level per tag value plus "missing") and gradient-boosted trees on slope and the covariates. The terrain-adjusted hiking function keeps Tobler's form but lets peak speed and slope sensitivity depend on one DEM-derived number, the steppiness z:

$$\ln v = a_0 + a_1 z - (k_0 + k_1 z)\,|s + s_0|, \qquad (1)$$

with z the SD of the 5 m pitches along the path divided by 0.10.

Predictive distributions and tests. Every model predicts its point estimate plus the quantiles of its residuals on held-out outings of the training data (training outings split 70/15/15 into fitting, early stopping and calibration), so calibration has to transfer too. Splits never separate a segment from its hike: 5 folds of outings (new hikers), 5 folds of 10 km UTM blocks (new places) and, the headline, leave one region out (17 folds); stricter checks hold out whole zones or half the continent. We report the continuous ranked probability score (CRPS) [17] of $\ln v$ as skill against the refitted slope curve, $1 - \mathrm{CRPS}/\mathrm{CRPS}_{\mathrm{slope}}$, with 95% intervals that resample outings (whole regions for unseen regions); the coverage of 90% intervals; and the median error in moving time over whole hikes after their first kilometre. A reference shows how much is the hiker: the median residual of a hike's first kilometre, shrunk by a factor fitted on calibration outings (0.61), applied to the rest of the hike. Asymmetry is measured on out-and-back pairs, the same outing on the same OSM way in both directions (mean grade at least 3%), as the ratio of speed down to speed up.

TABLE I
CRPS SKILL (%) VS THE REFITTED SLOPE CURVE. ADDITIVE / TREES; 95% INTERVALS FOR UNSEEN REGIONS RESAMPLE WHOLE REGIONS.

Model	New hikers	New 10 km	**Unseen region**
Tobler 1993	−12.8	−12.7	−12.5
Campbell 2022	−0.9	−0.9	−0.5
STRIDE 2024 (refit)	−17.2 (+0.2)	−17.3 (+0.2)	−16.8 (+0.2)
OSM tags	+1.3 / +1.5	+1.1 / +1.5	+0.9 / +0.4
30 m vegetation	+0.0 / +0.1	−0.0 / −0.1	−0.1 / −0.8
Embeddings	+1.4 / +2.6	+1.4 / +1.9	+1.1 / +1.1
Terrain	+2.2 / +3.2	+2.1 / +2.8	**+2.1 / +2.8**
Terrain + emb. + tags	+3.7 / +5.1	+3.8 / +5.1	+3.0 / +3.9
Eq. (1), one input	+1.0	+1.0	+1.0

IV. RESULTS

[TODO: R1 (Fig. 2a, Table I): Tobler fails everywhere (−12.5%). Walking speed peaks near 4.2 km/h on level to gently descending trail, not 6: Tobler is ∼20% too fast within ±10% grade and ∼18% too slow on climbs steeper than 25%. Campbell's published median curve matches near level but is ∼20% too fast on climbs steeper than 25%; it ties our refit on average, ranging by region from −3.8% (White Mountains) to +5.8% (Shenandoah).]

[TODO: R1b, asymmetry (Fig. 2b): on 2,763 out-and-back pairs (same hiker, same trail; 1,145 outings) descending is 8% faster than climbing at 3–10% grade, 22–23% above 20%; Tobler implies 42% at every grade above 5%. Our pooled curve implies nearly the same ratios as the pairs, so its asymmetry is not a sampling artifact; Campbell's curve implies less (2–16%).]

[TODO: R2: tags do not transfer: +1.3% within regions, +0.9% / +0.4% in unseen ones.]

[TODO: R3: terrain transfers: +2.1% / +2.8% in unseen regions, beating tags in 16 of 17 (Fig. 4); terrain minus tags +0.7 to +2.2, above zero in every bootstrap draw. About 60% of the gain is uneven grade inside the segment. Not a region proxy (adds ∼1.9 points over region offsets), not a hiker proxy. 1 m lidar ≈ 5 m; STRIDE's roughness adds nothing on trails; 30 m vegetation adds nothing; embeddings help within regions but transfer less reliably.]

[TODO: R3 robustness: the gain holds for newer hikes (trained up to 2019, tested on 2020 and later: +2.2% / +2.9%; tags +0.9%) and at every GPS recording rate (unseen regions, traces every ≤2 s, 2–10 s, >10 s: +1.7, +2.4, +1.9%). By slope it is largest on descents and level trail (+2.3 to +2.8%) and smallest on steep climbs (+0.8%). It holds with

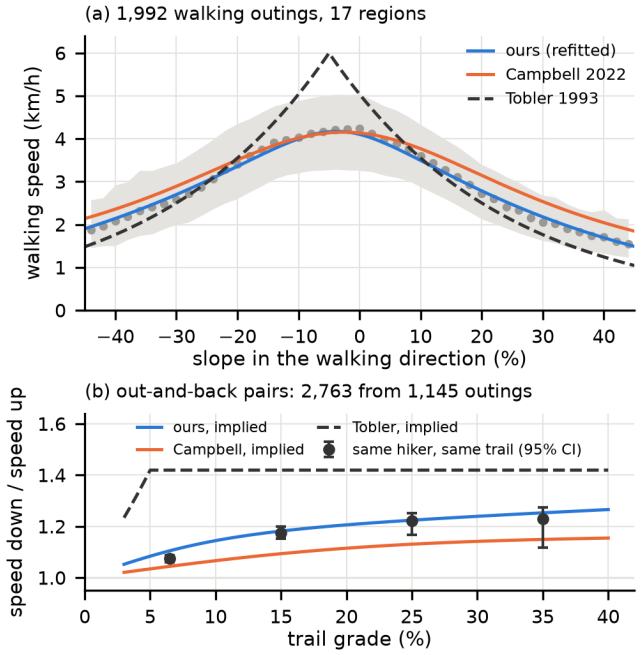


Fig. 2. (a) Walking speed vs slope in the walking direction, 17 regions: median per 2% bin (dots) and the middle 50% of segments (gray), our refitted curve, Campbell et al.'s published median curve [3] and Tobler's function [1]. (b) Speed down / speed up on out-and-back pairs (same hiker, same trail; grade bins 3–10, 10–20, 20–30, ≥30%), 95% intervals resampling outings, and the ratio each curve implies.

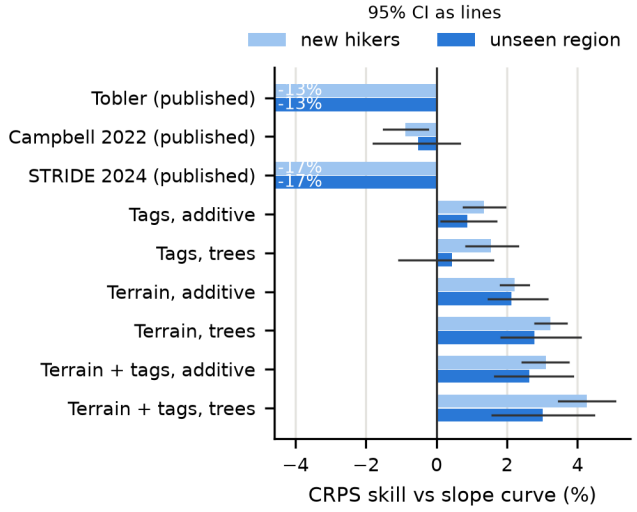


Fig. 3. CRPS skill vs the refitted slope curve for new hikers in known regions (light) and in regions the model never saw (dark), with 95% intervals.

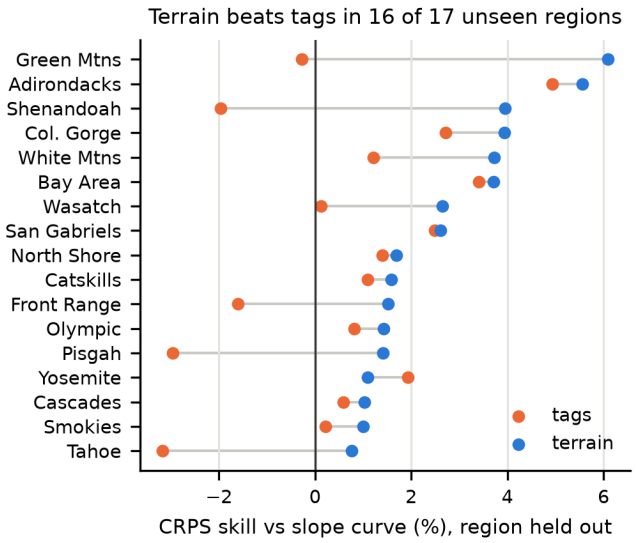


Fig. 4. Per held-out region: terrain vs OSM tags (additive models).

no neighbouring region in training: a whole zone held out (5 zones) +1.9% / +2.7%, trained on the West and tested on the East and back +2.6% / +3.3%; tags +0.7% / −0.1% and +1.5% / −0.3%; terrain minus tags above zero in every draw.]

[TODO: R4: the hiker dominates whole-hike time (Fig. 5): median error 13.6% (slope) → 13.0% (+ terrain); the hiker's first km brings it to 10.0%, and 9.6% with terrain.]

[TODO: R5, travel mode (plan claim 5): one curve times a per-mode multiplier predicts held-out runners as well as their own curve (×2.5; +0.1%), but not bikes and skis: their own curve is 3.8% better, and a constant ×3.2 is 15% too slow downhill and 35% too fast uphill.]

V. Discussion and Conclusion

[TODO: Why tags do not transfer (tested, decision #105): not because tag values mean different things. Within the same hike, harder sac_scale is slower in every region where it is used (−5% mountain_hiking vs hiking, 10 of 10 regions; −11% for harder ones, 8 of 8), and paved is faster (10 of 10). What does not carry over is whether a tag is there: sac_scale is on 1–96% of walked trail by region, and an untagged trail is faster than a tagged one in some regions and slower in others (−3.5 to +5.1% within a hike). So tags can only inform part of each region's trails, and how much differs by region; boosted trees on tags also lose most in transfer (+1.5% → +0.4%), as if fitting region-specific tag combinations. Elevation-derived terrain exists on every trail, and its effect has the same sign in all 17 regions (−12%, steppiest vs smoothest fifth).]

[TODO: What terrain captures: uneven grade within the segment (rock steps, short pitches) that one average slope hides, plus difficulty beyond grade. It matters most going down and on the level, where footing sets the pace, and least on steep climbs, where effort does. STRIDE's micro-roughness, calibrated off-trail, over-penalizes trails.]

[TODO: Practical output: Eq. (1) runs anywhere a DEM exists (Fig. 6); travel-mode multipliers fail for bikes (ratio ∼4× on descents, ∼1.6× on climbs).]

[TODO: Limits: OSM uploaders are not a random sample of hikers; sparse traces read slow, a measurement effect:

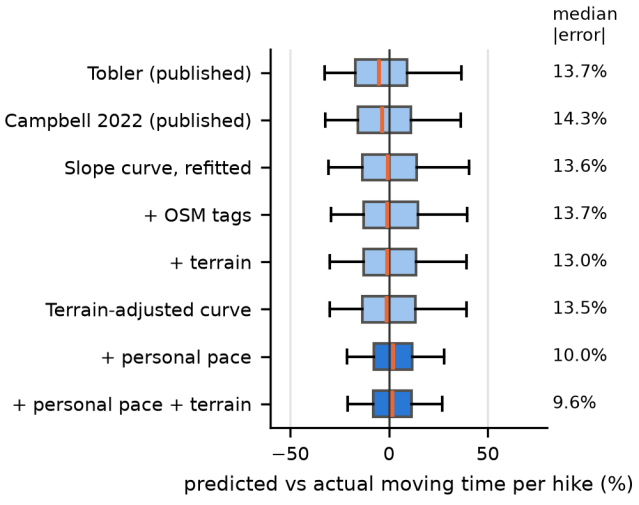


Fig. 5. Error in predicted moving time per hike, 1,354 hikes in held-out regions.

thinning dense walks to 30 s between points lowers measured speed by 11% (straight lines cut corners; short stops count as moving), and the curve without traces logged every >10 s peaks 3% higher (hiking functions depend on the sampling resolution [18]); per-segment effects are a few percent because GPS noise and the hiker dominate; weather [19] slows hikers on hot days (-3% above 25°C) and with snow on the ground (-4%), as in Hong Kong [5], but neither weather nor date and time of day improve predictions in new regions (weather -0.0%; with terrain its gain drops to $+1.9\%$), because the hiker's day-level pace already holds it; lidar understory density not tested (point-cloud cost).]

[TODO: Conclusion: three sentences. Data release (segments without user names) if the mentors agree.]

References

[1] W. Tobler, "Three presentations on geographical analysis and modeling," National Center for Geographic Information and Analysis (NCGIA), Tech. Rep. 93-1, 1993.

[2] M. J. Campbell, P. E. Dennison, B. W. Butler, and W. G. Page, "Using crowdsourced fitness tracker data to model the relationship between slope and travel rates," *Applied Geography*, vol. 106, pp. 93–107, 2019.

[3] M. J. Campbell, P. E. Dennison, and M. P. Thompson, "Predicting the variability in pedestrian travel rates and times using crowdsourced GPS data," *Computers, Environment and Urban Systems*, vol. 97, p. 101866, 2022.

[4] A. Wood, W. Mackaness, T. I. Simpson, and J. D. Armstrong, "Improved prediction of hiking speeds using a data driven approach," *PLOS ONE*, vol. 18, no. 12, p. e0295848, 2023.

[5] H. Huang, J. Y. Tsou, H. Zheng, R. Ma, H. Shi, and B. Zhang, "Modeling hiking speed via interpretable machine learning: Uncovering nonlinear impacts of multidimensional environmental factors," *Applied Geography*, vol. 192, p. 104052, 2026.

[6] L. Šerić, M. Tavra, I. Racetin, and A. Ivanda, "Modeling walkability by remote sensing as latent walking speed extracted from multiple digital trail maps," *Journal of Spatial Information Science*, no. 25, pp. 67–88, 2022.

[7] M. Asako, Y. Tahara, A. Ohsuga, and Y. Sei, "Deep learning-based travel time estimation in hiking with consideration of individual walking ability," *Cybernetics and Information Technologies*, vol. 24, no. 4, pp. 3–21, 2024.

[8] M. J. Campbell, S. L. Cutler, and P. E. Dennison, "A singular, broadly-applicable model for estimating on- and off-path walking travel rates using airborne lidar data," *Scientific Reports*, vol. 14, no. 1, p. 21393, 2024.

[9] M. J. Campbell, P. E. Dennison, A. T. Hudak, L. M. Parham, and B. W. Butler, "Quantifying understory vegetation density using small-footprint airborne lidar," *Remote Sensing of Environment*, vol. 215, pp. 330–342, 2018.

[10] A. Hadzic, H. Blanton, W. Song, M. Chen, S. Workman, and N. Jacobs, "RasterNet: Modeling free-flow speed using LiDAR and overhead imagery," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2020, pp. 826–834.

[11] OpenStreetMap contributors, "OpenStreetMap," 2026, GPS traces and map data, ODbL; https://www.openstreetmap.org.

[12] P. Newson and J. Krumm, "Hidden Markov map matching through noise and sparseness," in *ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, 2009, pp. 336–343.

[13] L. J. Sugarbaker, D. F. Eldridge, A. L. Jason, V. Lukas, D. L. Saghy, J. M. Stoker, and D. R. Thunen, "Status of the 3D elevation program, 2015," U.S. Geological Survey, Open-File Report 2016-1196, 2017.

[14] J. Dewitz, "National land cover database (NLCD) 2021 products," 2023, data release.

[15] C. F. Brown *et al.*, "AlphaEarth Foundations: An embedding field model for accurate and efficient global mapping from sparse label data," 2025, arXiv:2507.22291.

[16] Google and Google DeepMind, "AlphaEarth Foundations satellite embedding dataset," 2025, annual 10 m embeddings, CC-BY 4.0; read from the Source Cooperative mirror tge-labs/aef.

[17] T. Gneiting and A. E. Raftery, "Strictly proper scoring rules, prediction, and estimation," *Journal of the American Statistical Association*, vol. 102, no. 477, pp. 359–378, 2007.

[18] A. Goodwin, M. Hammett, and M. Harris, "The application of Tobler's hiking function in data-driven traverse modelling for planetary exploration," *Acta Astronautica*, vol. 228, pp. 265–273, 2025.

[19] M. M. Thornton, R. Shrestha, Y. Wei, P. E. Thornton, S.-C. Kao, and B. E. Wilson, "Daymet: Daily surface weather data on a 1-km grid for north america, version 4 R1," 2022.

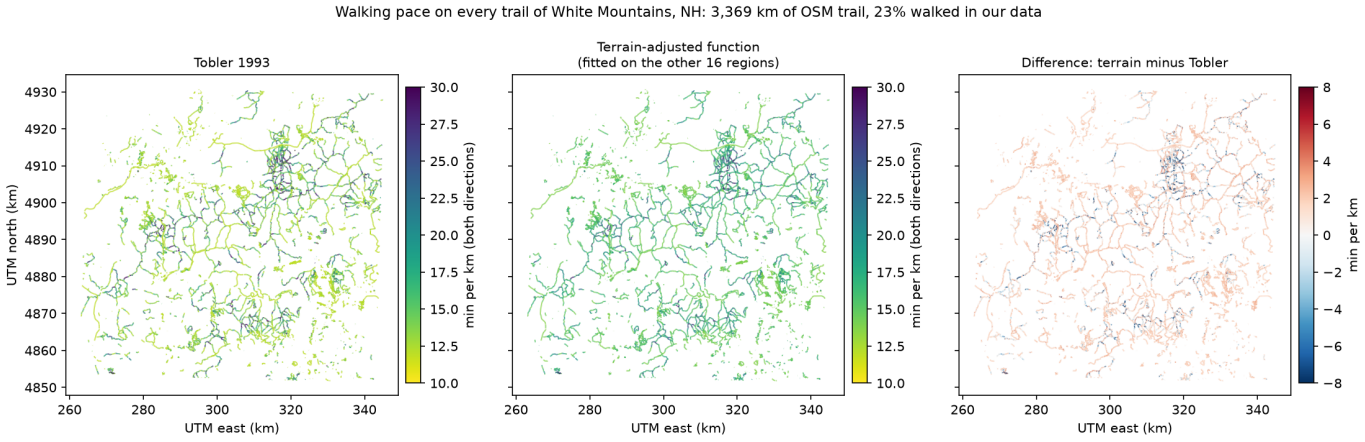


Fig. 6. Walking pace on every OSM trail of the White Mountains (3,369 km, 77% never walked in our data) from Eq. (1) fitted on the other 16 regions, vs Tobler. [TODO: Keep only if space allows; a two-panel crop may fit one column.]